

Introduction to probabilistic programming (with PyMC3)

A. Richards

02.08.2017

1 Introduction

2 PyMC3

3 Warm-up

4 PMF

5 Cool down

Probabilistic programming

A probabilistic programming language makes it easy to:

- 1 write out complex probability models
- 2 And subsequently solve these models automatically.

Generally this is accomplished by:

- 1 Random variables are handled as a **primitive**
- 2 Inference is handled behind the scenes
- 3 Memory and processor management is abstracted away

The pros and the cons

Why you might want to use probabilistic programming

- 1 **Customization** - We can create models that have built-in hypothesis tests
- 2 **Propagation of uncertainty** - There is a degree of belief associated prediction and estimation
- 3 **Intuition** - The models are essentially 'white-box' which provides insight into our data

Why you might **NOT** want use out probabilistic programming

- 1 **Deep dive** - Many of the online examples will assume a fairly deep understanding of statistics
- 2 **Overhead** - Computational overhead might make it difficult to be production ready
- 3 **Sometimes simple is enough** - The ability to customize models in almost a plug-n-play manner has to come with some cost.

Bayesian Inference

Degree of belief

You are a skilled programmer, but bugs still slip into your code. After a particularly difficult implementation of an algorithm, you decide to test your code on a trivial example. It passes. You test the code on a harder problem. It passes once again. And it passes the next, *even more difficult*, test too! You are [starting to believe](#) that there may be no bugs in this code...

Bayesian methods for hackers

This is a [nice intro to Bayesian thinking done on kdnuggets](#) (using PyMC3)

Some definitions

$$P(\theta|x) = \frac{P(x|\theta)P(\theta)}{P(x)} \quad (1)$$

- **prior** - $P(\theta)$ - one's beliefs about a quantity before presented with evidence
- **posterior** - $P(\theta|x)$ - probability of the parameters given the evidence
- **likelihood** - $P(x|\theta)$ - probability of the evidence given the parameters
- **normalizing constant** - $P(x)$
- $P(\theta)$: This big, complex code likely has a bug in it.
- $P(\theta|X)$: The code passed all X tests; there still might be a bug, but it is less likely now.

PyMC3

- Developed by John Salvatier, Thomas Wiecki, and Christopher Fonnesbeck (Salvatier et al., 2016)
- Comes with loads of good examples
- API is not backwards compatible with models specified in PyMC2

```
import pymc3 as pm

n,h,alpha,beta,niter = 100,61,2,2,1000

with pm.Model() as model: # context management
    # define priors and likelihood
    p = pm.Beta('p', alpha=alpha, beta=beta)
    y = pm.Binomial('y', n=n, p=p, observed=h)

    # inference
    start = pm.find_MAP() # initial state for MCMC
    step = pm.Metropolis() # Have a choice of samplers
    trace = pm.sample(niter, step, start)
```

Markov chain Monte Carlo (MCMC)

MCMC

- It is an family of algorithms for obtaining a sequence of random samples from a probability distribution for which direct sampling is difficult.
- The sequence can then be used to approximate the distribution
- It allows for inference on complex models

A particularly useful class of MCMC, known as Hamliltonian Monte Carlo, requires [gradient](#) information which is often not readily available so PyMC3 uses Theano to get around this problem. Another class of MCMC that has recently made this whole field a lot more interesting is the No-U-turn sampler (NUTS) because there are [self-tuning strategies](#) (Hoffman and Gelman, 2014).

One of the really nice things about probabilistic programming is that [you do not have to know how inference is performed](#), but it can be useful.

- MCMC for Dummies
- More on Hamliltonian MCMC (Not for dummies)

PyMC3 is an improvement over PyMC2

- Intuitive model specification syntax e.g.

$$x \sim N(0, 1) \text{ becomes } x = \text{Normal}(0, 1)$$

- Powerful sampling algorithms such as the [No U-Turn Sampler](#)
- [Variational inference](#): [ADVI](#) for fast approximate posterior estimation as well as [mini-batch](#) ADVI for large data sets.
- Relies on [Theano](#) which provides:
 - Numpy broadcasting and advanced indexing
 - Linear algebra operators
 - Computation optimization and dynamic C compilation
 - Simple extensibility
- Transparent support for [missing value imputation](#)

Getting started

We will be using Probabilistic Programming to perform automatic Bayesian inference on user-defined probabilistic models

- PyMC3 repo
- Getting started guide
- Bayesian methods for hackers

woot



coin flip example



Recommenders

The example

We will be walking through the [PyMC3 PMF example](#)





Where to go from here

Examples, examples, examples...

- [PyMC3 repo](#)
- [Getting started guide](#)
- [Bayesian methods for hackers](#)
- [Blog by Thomas Wiecki](#)
- [Doing Bayesian Data Analysis by John Kruschke](#)
- [Resource by Mark Dregan](#)

There is also [PyStan](#) ([Stan paper](#))

References I

Hoffman, M. D. and Gelman, A. (2014). The no-u-turn sampler: adaptively setting path lengths in hamiltonian monte carlo. *Journal of Machine Learning Research*, 15(1):1593–1623.

Salvatier, J., Wiecki, T. V., and Fonnesbeck, C. (2016). Probabilistic programming in python using pymc3. *PeerJ Computer Science*, 2:e55.